# THE DAILY BRIEF

OCTOBER 22, 2026

## AI & TECH

*What changed, and why it matters*

**OpenAI's newest models are now reachable through Google Cloud — and Microsoft is still in the picture.** OpenAI's current GPT 6 and 6.1 models — Astra, Sol, and Luna — will run on Google Cloud, appearing in Google's Model Garden alongside more than 200 other models and inside Antigravity, Google's agent-powered coding tool. Google says it is now the only large cloud provider offering all three leading closed models: Gemini, Anthropic's Claude, and OpenAI's. The models run entirely on Google's infrastructure, billed through the Cloud Marketplace, and are in private preview for select enterprise customers.

Why it matters: OpenAI still describes Microsoft as its primary cloud partner, so this is not a defection. It is a hedge. For a company whose economics depend on inference costs, having a second cloud is insurance — and for enterprises already committed to Google Cloud spending, switching to GPT models becomes a billing decision rather than an infrastructure project.

Read more: https://arstechnica.com/ai/2026/10/openais-latest-ai-models-are-coming-to-google-cloud-and-antigravity/

---

**A judge could only stop an AI model if 50 people might die.** Anthropic has been quietly pitching lawmakers a draft federal AI safety bill with a striking threshold. A model could be blocked only if it poses a foreseeable and imminent risk of more than $1 billion in damages, or of killing or seriously injuring more than 50 people — and only when that risk comes from biological weapons, the model evading its own controls, or a cyberattack. Only the Commerce Department could ask a federal judge; the draft explicitly bars any other federal agency from suspending a model. The harm thresholds closely mirror the catastrophic-risk language in California's Senate Bill 53.

Why it matters: safety bills are usually written by everyone except the companies they regulate. Anthropic writing one — with the trigger set above anything that has ever happened — is worth reading with that in mind. It is an attempt to define, in law, what counts as a genuinely dangerous model.

Read more: https://www.politico.com/news/2026/10/16/anthropic-ai-safety-01124654

---

**A coding model that runs on a laptop and helped build itself.** Recursive released Flint, a free coding model that runs on your own machine and was, the company says, built largely with the help of its own automated research system. It uses a mixture-of-experts design, activating just 2.5 billion of its 40 billion parameters at a time. In Recursive's tests it generated roughly 131 tokens per second on an M4 Pro MacBook Pro — about double the 58 to 69 tokens per second that Google's Gemma 4 and Nvidia's Nemotron 3.5 Lightning managed on the same machine. It is free to download under an Apache 2.0 license, with no cloud, no usage fees, and your code stays local.

Why it matters: this is self-improvement in a deliberately small wrapper — an AI research assistant overseen by humans that produced a model small enough to run locally. The honest limit is that no one has shown this loop running without a person in it.

Read more:
https://techcrunch.com/2026/10/21/recursive-launches-flint-a-laptop-coding-model-it-says-was-largely-built-by-ai/

---

*Generated by Marvin*