


THE DAILY BRIEF

Thursday, June 25, 2026



Markets & Finance

MAJOR MARKET MOVES AND FINANCIAL NEWS

JPMorgan Unveils \$50 Billion Buyback After Fed Stress Test Shows Banks Can Withstand \$708 Billion in Hypothetical Losses. The Federal Reserve released its annual stress test results Wednesday, and the headline was reassuring: all 32 large banks passed with room to spare, able to absorb \$708 billion in losses even under a severe recession scenario. JPMorgan immediately followed with a \$50 billion share repurchase authorization, signaling confidence in its balance sheet strength. Goldman Sachs raised its dividend on the same day. The results arrived at an interesting moment — the Fed simultaneously released an overhaul of capital rules that decouples capital requirements from stress test results, a move that banks have lobbied hard for. In practical terms, this means the stress test is now more of a transparency exercise than a binding constraint, which is why the buyback announcement matters more than the pass/fail headline. (CNBC)

The Wendy's Meme Stock Surge Shows Retail Traders Are Still Hunting for Viral Spikes. Wendy's shares surged more than 25% on Tuesday with trading halted at one point, then added another significant gain Wednesday as retail investors piled into what appears to be a social-media driven meme trade. The move was largely disconnected from company fundamentals — Wendy's management changes provided a thin narrative justification, but the magnitude of the move screams momentum chase rather than thesis-driven investing. This is the second notable meme stock episode this month, suggesting the appetite for short-duration viral trades among retail traders hasn't dimmed despite the broader market's more cautious tone. JPMorgan named two new co-presidents — Doug Petno and Troy Rohrbaugh — on Thursday, with longtime exec Marianne Lake exiting, in what looks like a planned succession reshuffle at the top of the nation's largest bank. (CNBC)

Kalshi CEO Hints at IPO Plans for the Prediction Market Platform. Tarek Mansour, CEO of Kalshi, told CNBC that the prediction market platform is thinking about going public but not this year. The comment is notable because Kalshi has been fighting the CFTC in court over whether prediction markets constitute regulated derivatives — the agency just sued Kentucky, the ninth state it's taken action against in its campaign to defend what it sees as its exclusive jurisdiction over event contracts. An IPO would represent a significant maturation event for the prediction market space, though the regulatory uncertainty around Kalshi's core business model makes the timing of any debut uncertain. (CNBC)

AI & Tech

LATEST IN ARTIFICIAL INTELLIGENCE, MODELS, AND TECHNOLOGY

Amazon Dumps \$13 Billion Into India's AI Infrastructure, Betting Big on the Country as a Next-Generation Compute Hub. Amazon announced a fresh \$13 billion commitment to AI infrastructure in India on Thursday, bringing its total planned investment in the country to jaw-dropping levels as U.S. tech giants race to establish footprint in markets that will host next-generation AI workloads. The investment targets data centers, cloud regions, and edge compute capacity across India, reflecting a view that India isn't just a services economy anymore — it's becoming a primary market for AI compute demand. This comes as Amazon, Microsoft, and Google have collectively pledged tens of billions into India's digital infrastructure over the past 18 months, a competitive dynamic that has drawn concern from India's own regulators about foreign dominance of critical compute infrastructure. (TechCrunch)

OpenAI and Broadcom Ship Jalapeño, a Custom Chip Built Specifically for Running Large Language Models at Scale. OpenAI and Broadcom unveiled their jointly developed inference chip, Jalapeño, on Wednesday, marking OpenAI's most concrete step yet toward reducing its dependence on Nvidia's GPU ecosystem. The chip is optimized for the specific computational patterns that LLM inference requires — matrix multiplications, attention computations, and memory bandwidth-heavy operations — rather than the general-purpose training workloads that Nvidia's H100 and B200 architectures are designed for. OpenAI's move mirrors what Google did with its TPU program years ago: build custom silicon purpose-built for your own workloads rather than paying premium prices for general-purpose chips. The announcement is a meaningful signal that the compute cost curve for AI is starting to bend as purpose-built silicon begins to compete with Nvidia's dominance. (OpenAI Blog, Ars Technica)

IBM Claims the World's First Sub-1 Nanometer Transistor Technology, a Potential Moonshot for Chip Performance. IBM's research division announced it has demonstrated working transistors below the 1 nanometer threshold — a milestone that semiconductor industry watchers have been projecting would arrive sometime in the early 2030s. The technology, called "nanosheet" or gate-all-around transistors, allows IBM to stack multiple sheets of silicon channel material in a way that dramatically increases drive current compared to today's FinFET designs. The practical implication, if the technology can be commercialized and manufactured at scale, is either a substantial jump in compute performance or a meaningful reduction in power consumption — or both. ASML's extreme ultraviolet lithography machines, which are required to etch these features, are already at the center of the U.S.-China chip war, and IBM's announcement will likely intensify the debate over which countries have access to which generations of chipmaking equipment. (Ars Technica)

AI Engineers Are the Most Resilient Job Category in Today's Labor Market, Defying the Layoff Narrative. New data from SignalFire, a venture firm that tracks hiring patterns across tech, shows that software engineers are actually growing as a share of total new hires — not shrinking despite the relentless coverage of AI replacing engineering roles. The finding cuts against the dominant layoff

narrative: while AI is demonstrably capable of handling specific coding tasks, code generation, and documentation, the incremental demand for engineers to build, test, and maintain AI systems themselves appears to be offsetting displacement at a systemic level. The pattern mirrors what happened when every "AI will kill coding" moment has arrived — demand for software talent has consistently grown, not shrunk, because AI lowers the cost of building and increases the surface area of what teams can attempt. (TechCrunch)

Google DeepMind Ships Computer Use in Gemini 3.5 Flash, Letting the Model Interact Directly with External Tools. DeepMind introduced computer use capabilities in its Gemini 3.5 Flash model Wednesday, a feature that allows the model to directly interact with external software tools, browsers, and APIs rather than just generating text responses. The move places Gemini in direct competition with OpenAI's operator-style agents and Anthropic's computer use features, all of which represent the industry's bet that the next frontier in AI utility is not just answering questions but taking actions on behalf of users. The computer use feature is a significant addition to Gemini's agentic toolkit and suggests the model is being positioned as a platform for building AI-powered workflows rather than just a chat interface. (DeepMind Blog)

AI Researchers Continue Their Exodus from Google, With Top Scientists Adler and Pritzel Joining Anthropic. Two more senior AI researchers — Jonas Adler and Alexander Pritzel — announced departures from Google this week to join Anthropic, adding to a growing list of high-profile exits that includes Noam Shazeer and John Jumper. The pattern reflects a broader redistribution of top AI talent away from large, bureaucratic tech companies and toward leaner AI-native organizations where researchers have more control over direction and compute resources. Anthropic has positioned itself as the destination for researchers who prioritize safety culture and research autonomy, and the continued departures from Google suggest that pitch is landing with a specific cohort of researchers who find Google's internal AI strategy and resource allocation frustrating. (TechCrunch)

Cerebras Stock Plunges After First Earnings Report as Investors Challenge the Chipmaker's Margin Narrative. The AI chipmaker's stock fell sharply Wednesday after its first earnings report since going public revealed a narrower gross margin outlook than investors had priced in. CEO Andrew Feldman said the market "misunderstood" the company's margin guidance, a response that tends to unnerve investors more than a straightforward restatement of fundamentals would. Cerebras occupies an interesting niche — its wafer-scale chips are physically enormous compared to Nvidia's offerings, which gives them advantages in certain memory-bandwidth-bound workloads like AI training, but the company has struggled to demonstrate the kind of broad adoption that would justify its valuation relative to Nvidia's ecosystem dominance. (TechCrunch)

Briefly

- › Europe Pushes Back on U.S. Chip War: European chip equipment maker ASML's CEO told TechCrunch that Washington's proposed MATCH Act would restrict China from buying older-generation deep ultraviolet lithography machines — gear first shipped about a decade ago — and that Europe has a different view on what level of semiconductor restriction makes sense.
 - › CFTC Sues Kentucky Over Prediction Markets: The derivatives regulator filed its ninth state-level lawsuit over prediction markets enforcement, escalating a legal battle that pits federal authority against states that have tried to restrict platforms like Kalshi and Polymarket.
 - › Vishal Sikka's New AI Services Venture: Former Infosys CEO Vishal Sikka unveiled a new startup backed by Mayfield and Aramco Ventures, aiming to bring AI-native approaches to enterprise IT services — a market that established players like Infosys and TCS have struggled to transform from legacy delivery models.
-

The Takeaway

The most consequential story this Thursday isn't any single development — it's the compounding evidence that the AI infrastructure buildout is accelerating on multiple fronts simultaneously. Amazon's \$13 billion India commitment, OpenAI's custom Jalapeño chip, IBM's sub-1nm transistor claim, and DeepMind's computer use launch all point in the same direction: the cost and capability curves for AI are moving fast, and the companies building the compute layer — not just the model layer — are where the durable competitive advantages are being carved out. For markets, JPMorgan's \$50B buyback signals that large-cap banks remain in robust shape post-stress test, while the Wendy's meme surge is a reminder that retail momentum trades haven't gone extinct. The Kalshi IPO hint is worth watching — if prediction markets ever get regulatory clarity, the access to retail capital markets would be a genuine structural shift for that industry.
