

THE DAILY BRIEF

Friday, June 26, 2026

AI & Tech

THE LATEST IN ARTIFICIAL INTELLIGENCE, MODELS, AND TECHNOLOGY

The White House Told OpenAI to Keep GPT 5.6 Off the Market — And OpenAI Listened. In a remarkable display of government pressure on AI development, the Trump administration has asked OpenAI to restrict its newest model to a small group of partners rather than releasing it broadly. The model, internally referred to as GPT 5.6, was set for a wider debut but was quietly scaled back after direct intervention from the White House over safety concerns. The episode underscores how deeply the AI race is now intertwined with federal policy — and how much leverage the government believes it holds over the most advanced labs. (TechCrunch)

Anthropic Says Alibaba Ran the Largest Attack Ever to Clone Claude — and Did It While Defying Trump. Anthropic has handed Congress what it calls new, confidential evidence of an audacious campaign: 25,000 fraudulent accounts generating 28.8 million exchanges with Claude between April and June, all orchestrated by operators affiliated with Alibaba and its AI lab, Alibaba Qwen. The targets were Claude's most valuable capabilities — agentic reasoning, software engineering, and long-horizon tasks. What's striking is the timing: the attacks occurred even as the Trump administration has been restricting Chinese access to frontier AI models like Anthropic's Mythos. Anthropic says Alibaba effectively used the restrictions as cover, attacking precisely the models it was supposed to be locked out of. (Ars Technica)

Anthropic's Claude Is Quietly Eating ChatGPT's Lunch in the Paid Consumer Market. ChatGPT still commands the overall lead in AI subscribers, but consumers who pay for AI are increasingly choosing Claude instead. New data from Anthropic shows a steady migration of paid users toward Claude — a trend that reflects something important about where the market is heading. Claude's strength appears to be depth: better reasoning, longer context windows, and a writing style that power users find more reliable for complex tasks. ChatGPT remains the default for casual users, but in the tier that actually pays monthly, Claude is gaining ground fast. (TechCrunch)

Patronus AI Lands \$50 Million to Build Digital Worlds That Stress-Test AI Agents. The agentic AI boom is creating a new bottleneck: how do you actually evaluate whether an AI agent is safe, reliable, and

working correctly when it operates across multiple steps and tools? Patronus AI, founded by former Meta researchers, is tackling exactly that problem — building simulated environments that put AI agents through their paces before they touch real systems. Investors are paying attention: the company just closed a \$50M round led by Greenfield Partners and Lightspeed. The pitch is essentially "AI red-teaming at scale," and demand is described as nearly insatiable. (TechCrunch)

General Intuition Has Raised \$320 Million to Train AI on Video Games and Deploy It in the Real World. The premise sounds like a pitch deck from 2016, but General Intuition is making a serious bet that millions of hours of human gameplay can teach AI something that text and code cannot: how humans actually behave in complex, adversarial environments. The company has raised a \$320M Series B at a \$2.3B valuation, with Khosla Ventures leading, to scale what it calls "action data" training. The idea is that world models trained on gameplay develop something closer to human intuition — an AI that can reason about physical and social environments the way a person does, not just pattern-match on tokens. (TechCrunch)

Databricks' Former AI Chief Says He Can Cut AI's Power Bill by 1,000x. The single biggest threat to AI's continued growth isn't model capability — it's electricity. Training and running large models is extraordinarily energy-intensive, and as demand scales, power availability is becoming a literal bottleneck on expansion. Now a former Databricks AI chief is going public with a solution he claims can reduce AI's power consumption by three orders of magnitude. The approach involves rethinking inference architecture from the ground up, moving away from the dense, wasteful computations that characterize current large-model deployments. No commercial product yet, but the claim alone has gotten the industry's attention. (TechCrunch)

OpenAI and Broadcom Unveil a Custom AI Chip Built for LLM Inference. OpenAI has been quietly working with Broadcom on a purpose-built inference chip, and the two companies just announced the result: Jalapeño. The chip is optimized specifically for running large language model inference — not training — which is where most of the compute cost actually accumulates in production deployments. This is a significant strategic move: OpenAI has been reliant on Nvidia for training, but for inference at scale, custom silicon offers better cost-per-token economics. Jalapeño is OpenAI's first concrete step toward controlling more of its own hardware stack. (OpenAI Blog)

Markets & Finance

MAJOR MARKET MOVES AND FINANCIAL NEWS

Polymarket Says Hackers Stole Users' Funds — And Is Now Refunding Them. The prediction market giant disclosed a security breach affecting user accounts, with funds apparently siphoned out through a third-party exploit. Polymarket says it is actively refunding affected users, though the company has not disclosed the total amount stolen or the specific vulnerability exploited. This is an awkward moment for a platform that has become a central nervous system for political and financial forecasting — its credibility rests entirely on being trustworthy enough to put real money on. A breach of this kind, even if fully compensated, raises questions about the robustness of its security architecture. (TechCrunch)

Apple Raised Mac and iPad Prices Overnight, Spared iPhone — For Now. Apple's memory and component costs have more than doubled relative to previous product cycles, and the company has finally passed that pain on to consumers — at least in the non-iPhone categories. MacBook Air, Mac Pro, and the iPad Pro all saw price increases, with the most popular configurations seeing bumps of \$100 to \$300. iPhone pricing was left unchanged, which is notable: Apple clearly doesn't want to risk further erosion in its highest-volume, highest-margin product category while Android competitors grow more aggressive on price. But the trajectory is clear — hardware is getting more expensive to make, and the era of \$999 entry-level iPads is over. (TechCrunch)

Microsoft Quietly Extended Windows 10's Life Support Another Year. Windows 10's official support death date was October 2025, which was supposed to trigger an exodus to Windows 11. That exodus hasn't happened fast enough for Microsoft: as of today, roughly a quarter of all PCs are still running Windows 10. So Microsoft quietly extended its Extended Security Updates program from an intended October 2026 cutoff to October 2027, giving stragglers another full year of patches. The underlying problem — Windows 11's hardware requirements are too demanding for a large chunk of the installed base — hasn't been solved. It's just been deferred. (Ars Technica)

Google Finance Finally Has Its Own Mobile App, and It's Entirely About the AI. Twenty years after Google Finance launched as a website, Google has released its first standalone Android app — and the iOS version isn't coming until later this year. The app is built around a complete AI overhaul of the Finance website, with Google's AI generating "key moments" that explain why specific stock prices moved. The timing is curious: a finance app without mobile coverage for half the smartphone market seems like an incomplete launch. But Google appears to have prioritized getting the AI features right over hitting both platforms simultaneously. (Ars Technica)

Briefly

- › FCC Chair Proposes Killing \$2B Program That Connects Schools to the Internet. FCC Chairman Brendan Carr is pushing to scale back or eliminate E-Rate, a program that provides telecom discounts for schools and libraries. His reasoning: students are getting too much screen time. Critics call it a radical overreach that would gut connectivity for the institutions that need it most. (Ars Technica)
 - › Trump Administration Proposes Removing Brake Pedal Requirement for Autonomous Vehicles — a Boon for Tesla. The National Highway Traffic Safety Administration is floating a rule change that would eliminate the requirement for traditional brake pedals in fully autonomous vehicles. Currently, even "self-driving" cars must have manual controls as a failsafe. Removing that mandate would streamline the regulatory path for vehicles like Tesla's Robotaxi, which has no steering wheel or pedals. (TechCrunch)
 - › Polestar Barred From Selling New EVs in the US From 2027 Onward. The Commerce Department has denied Polestar authorization to import vehicles from model year 2027, citing the connected car rule that targets automakers with Chinese corporate parentage. Polestar is owned by Geely, the Chinese conglomerate that also controls Volvo. Unlike Volvo, which received a separate exemption, Polestar gets no such reprieve. The Polestar 3 and 4 can still be sold from existing inventory. (Ars Technica)
 - › Rocket Lab Executed a US Space Force Mission 16 Hours After Getting the Call. The company launched its Victus Haze mission — preparing, launching, and initial deployment — in just 16 hours and 42 minutes from notice to launch. The previous record was over 27 hours. The mission was barely announced in advance; the only public signal was a NOTAM (Notice to Airmen) telling pilots and sailors to clear the flight path. No livestream. No fanfare. Just execution. (Ars Technica)
-

The Takeaway

The dominant theme of this Friday is the growing tension between AI capability and AI control — and it's playing out on multiple fronts simultaneously. The White House pressuring OpenAI to restrict GPT 5.6, Anthropic exposing Alibaba's audacious model-cloning operation, the FCC moving to cut school internet funding over screen time concerns: each is a different flavor of the same underlying story. AI is no longer just a technology story. It's a policy story, a geopolitics story, a trade story, and an infrastructure story — sometimes all at once. For investors, the signal is clear: the bottlenecks on AI aren't going to be compute or data. They're going to be regulatory, political, and physical. Pay attention to who controls the compute stack, who controls the data pipelines, and who controls the regulatory levers. Those are the moats now.
